

Changing a Networked Mind: Paradigm Dynamics as Structure Learning over a Hidden World

Jonas Hallgren

David Hyland

Mahault Albarracin*

Draft — June 11, 2026

Abstract

A scientist is not handed the variables of the world, nor the dependencies among them: the structure of reality is hidden, and a community must *learn* that structure even as it uses it. We model a paradigm as a directed network of commitments and a scientific community as a population of bounded active-inference agents performing *Bayesian structure learning* on that network—expanding it by proposing new commitments and reducing it by pruning the dependencies the data do not support. One operator organizes the formalism: the propagation operator of the network carries a *conservatism* field (the cost of revising a commitment) and a *conviction* field (the value a bounded agent attaches to holding it), the same operator applied to two different sources. Running the model through the full Kuhnian cycle—normal science, anomaly, crisis, reduction, discovery—in a 2,056-run sweep yields three findings. The lock-in boundary sits at *very* weak social coupling: conviction-driven gating of disconfirming evidence cannot defeat connection, and only isolation protects a superseded paradigm. Discovery of a genuinely unconceived dimension follows a Poisson waiting-time law, but a *staggered* discovery never survives the communal prior—a concept persists only when the community wakes it near-simultaneously, a critical mass in *time*. And one clean negative: a conservatism-gated revision rate does *not* recover the predicted belt-first/core-last staircase, which localizes the missing ingredient to data-dependent structure learning.

1 Introduction: changing a frame, not a credence

To change one’s mind about a fact is to update a belief; to change a paradigm is to change the *frame* in which beliefs are posed. Such a frame is not a credence but a *structure*: a network of conditional dependencies through which information propagates. A community that only revised its credences, leaving the network fixed, could not change frame at all. This is the phenomenon Kuhn named but could not run forward (Kuhn, 1962): nothing in *The Structure of Scientific Revolutions* tells us, mechanistically, why a community holds a structure of commitments long past the first disconfirming evidence, nor what finally makes the holding give way.

When the chemical community released phlogiston it did not revise a single credence. It rewrote a network of commitments in which combustion, calcination, respiration, and the theory of acids were all load-bearing on a hub that had to go: late-phlogistic theorists proposed cascading revisions that proved mutually incompatible (Blumenthal and Ladyman, 2017; Holmes, 2000), defending less a hypothesis about combustion than a constellation of entangled commitments (Boantza and Gal, 2011)—phlogiston was load-bearing for too much.

The dominant formalism for such communities, network epistemology, cannot represent this case. Its agents hold a scalar credence in a fixed proposition and update on a trust graph (Zollman,

*Author order and affiliations provisional. Draft for IWAI 2026.

2010; O’Connor and Weatherall, 2019; Zollman, 2013), a programme that has produced striking results on transient diversity, polarization, and the division of cognitive labour (Jonard et al., 2025; Michelini et al., 2023; Kitcher, 1990; Strevens, 2003)—but the phlogiston puzzle is not a failure to converge on a number. The commitments were *entangled*: revising any one forced a re-reading of the others, and the disagreement was about which entanglements to keep. Thagard’s ECHO (Thagard, 1989, 1990) captures the structural intuition as undirected coherence; it lacks the directed dependencies, the carry-over cost of mind-change, and the expand/reduce asymmetry we derive below. Active-inference accounts of epistemic communities supply an endogenous precision-based lock-in mechanism (Albarracin et al., 2022) and warn that naive aggregation degrades collective precision (Friston et al., 2024a; Waade et al., 2025); the variational cost of mind-change has been formalized for a single agent as a motivated trade between revision cost and the utility of a held belief (Hyland and Albarracin, 2025). What none of these does is treat *belief structure itself* as the object of inference: they retain a fixed menu of hypotheses, and a scalar credence over a menu cannot say how revising one commitment forces a re-reading of the rest, nor how a community arrives at a *new* way of carving up the world. When only a number moves on the trust graph, a community can converge on it, herd on it, or split over it—but its members can never come to hold different *kinds* of world, only different values on a shared one. The question, made concrete in §3.5, is therefore what moves: this paper lets a piece of *structure* travel, so that one agent can come to represent a commitment its peers do not.

We make structure the object. Agents carry a directed network of commitments, learn its structure by Bayesian model expansion and reduction, and are *bounded*: they sample a complex world through narrow channels, value some conclusions over others, and judge what they receive against their own model. Section 3 builds the model and states the simulation loop explicitly; Section 4 reports a single-agent dark-energy net and a 2,056-run population sweep through the full Kuhnian cycle on the phlogiston world.

2 Why paradigms are networks under structure learning

The structure of reality is hidden. A scientist confronting a new domain is not given the right variables, and is certainly not given the dependencies among them: she does not know whether two regularities share a hidden common cause, are lawfully coupled, or merely coincide. For an active-inference reader the move is natural: just as agents infer hidden states, they must infer the hidden *dependencies among* those states. The cognitive-science literature on theory acquisition reaches the same conclusion—conceptual development as hierarchical Bayesian program induction, a stochastic search over structures (Ullman and Tenenbaum, 2020; Kemp and Tenenbaum, 2008; Ullman et al., 2012).

To explore is to expand; to do normal science is to reduce. Network-epistemology models keep their communities from collapsing onto the first attractive theory with a hand-tuned exploration term—an ϵ , a temperature, a switching cost (Zollman, 2010; O’Connor and Weatherall, 2019). Once structure is the object of inference, the *content* of exploration is no longer a tuned bonus but the epistemic value of expanding the model to entertain a commitment not yet in it (Friston et al., 2015, 2017); what stays free is only the *rate* at which proposals arrive. The complementary operation is Bayesian model reduction: pruning structure the data do not support, in closed form, at the slowest timescale of the learning hierarchy (Friston et al., 2018; Smith et al., 2022b). Existing applications of this machinery run on single agents (Smith et al., 2022b; Priorelli and Stoianov, 2023) or define “group” as between-subject neuroimaging (Friston et al., 2016); treating a paradigm as the object a

community jointly reduces and expands is the open move this paper makes.

Two fields, one operator. The single non-obvious step the rest of the paper turns on: once a paradigm is a network, one operation—propagating influence from a commitment to everything downstream—generates *two* fields on the same graph, depending on what is propagated. Propagate sheer dependence and we get *conservatism*: a commitment is costly to revise because much depends on it. Propagate *value* and we get *conviction*: the worth a bounded agent attaches to holding a belief, carried through the same dependencies. The two are independent—a commitment can be central yet value-neutral, cheap yet cherished (Fig. 1). That attitudes live on a network is the empirical posture of the Causal Attitude Network model (Dalege et al., 2016); we place it on the same *directed* graph that carries the cost of revision.




Figure 1: Two fields on one paradigm (schematic). Both panels carry the *same* network—a densely coupled core $\{A, B, C, D\}$ and a sparse belt $\{b_1, \dots, b_6\}$ in contact with observation. a Revision cost is the carry-over of mind-change, $\kappa = \mathbf{T1}$ (Eq. 1): a node’s cost is the precision-weighted mass of its descendants, a clean radial gradient monotone in centrality. b Intrinsic utility is signed—beliefs one wants true (green), wants false (red), or is neutral toward (white)—and the conviction field $U = \mathbf{T}u$ does not align with κ : C is central yet value-neutral; b_2 is cheap yet strongly held; b_4 , which feeds C , is a conclusion the community wants false. The two fields are $\mathbf{T1}$ and $\mathbf{T}u$, the same operator on two sources, and they are independent vectors. Appendix B shows this decoupling computed on a real dark-energy net.

3 The model

We build the object up from the single agent and then state the population dynamics as an explicit per-round loop (§3.5).

3.1 A paradigm network and its operator

An agent represents a paradigm as a directed network $G = (V, E, \Pi)$: nodes carry theoretical commitments x_v , edges carry inferential dependencies, and each edge bears a coupling precision π_{uv} measuring how tightly v is bound to u . Internal coherence is a precision-weighted graphical model, $p(x | G) \propto \prod_{(u \rightarrow v) \in E} \psi_{uv}(x_u, x_v)^{\pi_{uv}}$. The dense, high-precision interior is the *core*; the sparse periphery in contact with observation is the *belt*, each belt node reading the world through a sensory channel of precision ρ_k . The core carries no direct sensory channel: an interior commitment is revised only by error propagated inward from the belt, which is what makes it slow. We work throughout in a *linear-Gaussian* regime, which puts every quantity below in closed form (Appendix A); agents are *honest*—no one misreports, so lock-in is never a dishonesty artifact—and *bounded*, minimizing variational free energy step by step (Da Costa et al., 2020) rather than planning over policies.

Revising a commitment is not a local act: changing x_v forces everything that depends on it to re-equilibrate. The whole cascade is generated by one object. From the adjacency A , the directed *propagation operator* is $\mathbf{T} = (I - A)^{-1}$, whose entry $\mathbf{T}_{iw} = \pi_{i \rightsquigarrow w}$ is the total precision-weighted coupling propagated from i to a descendant w along all directed paths (Appendix A). Conservatism is $\mathbf{T}$ applied to the unit source,

$$\kappa = \mathbf{T} \mathbf{1}, \quad \kappa_i = 1 + \sum_{w \in \text{desc}(i)} \pi_{i \rightsquigarrow w}, \quad (1)$$

the precision-weighted mass of a node’s descendants: a core node is costly to revise because much depends on it, a belt node cheap because little does. Conservatism is thus *derived* from connectivity, not assigned by hand.

The second field is value. Scientists are bounded agents who value some conclusions over others—a result on which a career rests is held more tightly than its evidence warrants (Kitcher, 1990; Strevens, 2003). Following the variational-cost account of belief revision (Hyland and Albarracin, 2025), we attach to each node an intrinsic utility u_i and let it propagate through the *same* operator: a commitment one wants true lends value to its dependents, so the conviction field is

$$U = \mathbf{T} u, \quad U_i = u_i + \sum_{w \in \text{desc}(i)} \pi_{i \rightsquigarrow w} u_w. \quad (2)$$

$\mathbf{T} \mathbf{1}$ and $\mathbf{T} u$ are linearly independent unless $u \propto \mathbf{1}$: the two fields need not align (Fig. 1), and Appendix B exhibits the misalignment computed on a real net.

3.2 Motivated revision and the lock-in parameter

The bounded agent revises its belief $q(x)$ over the whole network by trading fidelity to the honest posterior against the value of where the commitments land,

$$\mathcal{F}[q] = D_{\text{KL}}[q(x) \| p(x | o, G)] - \lambda \mathbb{E}_q[U^\top x], \quad \implies \quad q_\lambda(x) \propto p(x | o, G) e^{\lambda U^\top x}, \quad (3)$$

with λ a single temperature governing how far value may bend the evidence. In the linear-Gaussian regime the tilt is a pure mean shift, displacing the commitments by $\lambda \Sigma U$ and leaving the covariance untouched (Appendix A). What opposes the displacement is evidence, also through the mean: a disconfirming channel k with sensory precision ρ_k and read-out H_k pulls the mean back with gain $\propto \rho_k H_k \Sigma U$. As conviction drives a disconfirming channel toward $\rho_k \rightarrow 0$ that gain vanishes: the

channel can no longer move the mean, and the value-displaced commitment stands unrefuted. The fraction of disconfirming sensory precision so silenced,

$$\gamma = \frac{\sum_{k \in \text{dis}} \rho_k \mathbb{I}[\rho_k \text{ gated by the core}]}{\sum_{k \in \text{dis}} \rho_k}, \quad (4)$$

is the *lock-in parameter*. In the experiments γ is endogenous: each round the weight on a disconfirming row is a saturating function of the conviction field projected onto it, $w_k = \exp(-g |U \cdot H_k|)$, with g the conviction-to-censorship coupling—no channel is gated by hand.

3.3 Structure learning: reduction and expansion

Structure learning revises G itself through two asymmetric operations (Friston et al., 2017; Smith et al., 2020), and the asymmetry is the formal counterpart of normal versus revolutionary science.

Reduction is the exact half. Each edge is parameterized by a prior precision α_e ($\alpha_e \rightarrow \infty$ pins the edge’s parameter to zero; $\alpha_e \rightarrow 0$ frees it), the automatic-relevance-determination parameterization of sparse Bayesian learning (MacKay, 1992; Tipping, 2001). Bayesian model reduction (Friston and Penny, 2011; Friston et al., 2018) scores any pruning in closed form from the already-inverted posterior,

$$\frac{p(o | \tilde{M})}{p(o | M)} = \mathbb{E}_{p(\theta|o)} \left[\frac{\tilde{p}(\theta)}{p(\theta)} \right], \quad (5)$$

the Savage–Dickey ratio in general form: one inversion scores every candidate reduced structure without re-touching the data.

Expansion reaches what reduction cannot: a commitment no one has represented. Following Friston et al. (2017), the agent entertains a new node by wiring it tentatively to existing nodes with free priors ($\alpha_e \rightarrow 0$) and letting reduction prune the proposal back to the couplings the data support (Fig. 2). Expansion is not closed-form—a new node may enlarge the likelihood’s support and requires a real inversion—and this cost asymmetry is why a field reorganizes its periphery continually and its centre only rarely. The agent does not plan *when* to expand: candidates arrive at a Poisson rate and are scored by ΔF . It must be a rate, because one cannot plan toward a dimension one has not yet conceived—Stanford’s problem of unconceived alternatives (Stanford, 2006), to which we return.

3.4 Forgetting

Read once from fixed precisions, $\kappa = \mathbf{T}\mathbf{1}$ can only grow as learning adds edges: a paradigm with no forgetting is a ratchet, and lock-in under it is automatic rather than earned. We import the standard remedy unmodified: accumulated precision relaxes toward the prior between updates, $\Pi \leftarrow \Pi_0 + \omega(\Pi - \Pi_0)$ with forgetting factor $\omega \in (0, 1]$ (Friston et al., 2017b; Smith et al., 2022), so a deposit k steps old carries weight ω^k . The factor is the agent’s prior on the world’s structural drift—one forgets at the rate one believes the structure is moving (Mathys et al., 2011). The operator inherits the time index, $\mathbf{T}_t = (I - A_t)^{-1}$, and a once-central node whose couplings go unreinforced relaxes back toward cheapness. With forgetting, the value tilt saturates at $\lambda U / (1 - \omega)$ and competes with the also-saturated per-step evidence, so conviction holds only above a threshold: in earlier population runs the entrenched bloc re-tracked the truth for $\lambda \leq 6$ and stayed locked for $\lambda \geq 15$ —motivated persistence turned from a posited bias into a threshold one must clear.




Figure 2: Structure learning as expansion and reduction. a To entertain a commitment not yet in the model, the agent proposes a new node N wired tentatively to every existing commitment with free priors ($\alpha_e \rightarrow 0$, dashed). b Bayesian model reduction prunes the proposal in closed form, keeping only the couplings the data support. Expansion costs a real inversion; pruning is the cheap closed-form half.

3.5 The population, and the simulation at a glance

A population of N agents inhabits a shared hidden world, each carrying its own paradigm network with its own coupling precisions, conviction field $U = \mathbf{T}u$, and derived conservatism $\kappa = \mathbf{T}\mathbf{1}$. Agents are coupled on a trust graph: communities are blocks of a stochastic block model, with dense intra-community edges and an inter-community coupling weight (*inter*) that we sweep from total isolation to dense contact. Beliefs and accumulated precisions are fused across trust edges in information form—a precision-weighted pooling of posteriors in which an agent’s own evidence and its neighbours’ carry weight in proportion to trust—and structural edits spread the same way: an agent whose neighbour has awakened a new node receives the proposal and scores it against its own model by Eq. (5). Figure 3 draws the picture: what travels on the trust graph is not a number but a piece of structure, evaluated against wiring the sender does not see.

Concretely, the experiments of §4.1 instantiate the phlogiston world: a hub-structured paradigm net in which a single core commitment (phlogiston) is load-bearing for combustion, calcination and respiration; at a fixed step t_{shift} the world’s regime flips (the calx is heavier—the gravimetric anomaly), and the truth thereafter contains a dimension (oxygen) that is *not in any agent’s network*. One round of the simulation, for every agent:

1. **Observe.** The world emits observations at the belt; channel weights on disconfirming rows are set endogenously by the conviction gate, $w_k = \exp(-g |U \cdot H_k|)$ (§3.2), scaled by the agent’s gate strength s .
2. **Infer.** The agent updates its (value-tilted, Eq. 3) posterior over x by exact linear-Gaussian inference and accumulates precision.
3. **Crisis check (reduction).** The closed-form prune evidence ΔF (Eq. 5) on the hub’s edges is compared against the agent’s conviction protection $\lambda_i v_e$; a crossing is a *crisis*, and the agent applies the reduction—the hub is released.




Figure 3: The population model. Right: the community is a trust graph, every circle an agent. Left, magnified: each agent privately holds a paradigm network $G_i = (V, E, \Pi)$ —commitments coupled by precisions π_{uv} , a belt reading the hidden world through sensory channels of precision ρ_k —from which its conservatism $\kappa = \mathbf{T}\mathbf{1}$ and conviction $U = \mathbf{T}u$ are derived. What passes along trust channels (amber) is a *structural proposal*: a piece of paradigm, e.g. a newly awakened node, which each receiver scores against its own network by ΔF (Eq. 5); posteriors fuse precision-weighted along the same channels. The trust graph is fixed at the start (a stochastic block model): channels are dense and strong inside a community, and a single weak weight *inter*—the main dial of §4—connects communities, so what reaches an agent from outside is its neighbours’ evidence and proposals, scaled by trust.

4. **Discovery check (expansion).** After reduction, the residual the old paradigm absorbed becomes visible as a coherent floor in prediction error; when the floor clears the expansion trigger (and a Poisson proposal has arrived, rate r), the agent wakes the unconceived node and wires it in (§3.3).
5. **Fuse.** Posteriors and structural edits are pooled across the trust graph, weighted by the intra/inter coupling.
6. **Forget.** Accumulated precision decays by ω (§3.4).

Crisis, revolution and discovery are thus not stipulated events but threshold crossings of the model’s own evidence. Table 1 lists the dials swept; everything else is held fixed, and all code is seeded and ships its own null and positive controls.

Dial	Symbol	What it controls
Social coupling	<i>inter</i>	inter-community trust weight (0 = isolation)
Vanguard fraction	—	fraction of agents in the open (low- λ) community
Conviction	λ_d	dogmatic community’s value tilt / prune protection
Gate scaling	s	strength of the conviction gate on disconfirming channels
Proposal rate	r (Poisson)	arrival rate of structural proposals
Expansion trigger	—	residual floor required to wake a new node
Forgetting	$\omega \in (0, 1]$	decay of accumulated evidence toward the prior
Observation noise	σ_o	sensory noise at the belt

Table 1: The dials swept in the experiments (2,056 runs); all other quantities are held fixed.

4 Results

All results are from the population of §3.5 run through the full Kuhnian cycle on the phlogiston world, swept overnight over the dials of Table 1 (2,056 healthy runs). The single-agent *dark-energy* net that realises the two fields of Fig. 1 on a worked example—and recovers the Ω_m – Λ degeneracy as a Schur complement, a quantity a scalar credence has no room to express—is reported in Appendix B. The nets are hand-wired to expose the geometry, not fitted to data: nothing here is evidence about chemistry or cosmology, only about how a structured, bounded, valued mind changes. We close with limitations and the model’s one clean negative (§4.2).

4.1 The full cycle, endogenously: the phlogiston sweep

In the population, the whole Kuhnian arc—normal science, the gravimetric anomaly, crisis, reduction, discovery of oxygen, new normal—unfolds as successive threshold crossings of the agents’ own evidence (§3.5), with no phase stipulated. The sweep then asks where each beat of the arc lives in the dial space, and four findings organize it.

Even a trace of outside contact ends the old paradigm; only total isolation protects it. The first question the sweep answers: how much contact with the rest of the world can a dogmatic community tolerate and still keep phlogiston? Almost none. The setup pits a dogmatic community—high conviction, which makes its agents gate their own disconfirming channels (§3.2)—against an open vanguard already tracking the anomaly, and dials the trust weight between the two communities (*inter*) from 0 (sealed off) up to dense contact. Sealed off, the protection works: under strong gating ($s = 10$) only 19% of the dogmatic community ever converts to oxygen. But the seal has to be perfect. At *inter* = 0.005—trust toward outsiders two hundred times weaker than trust within the community—more than 93% convert, no matter how strong the gate, and making the agents twice as dogmatic ($\lambda_d : 0.35 \rightarrow 0.8$) barely shifts the picture (Fig. 4). The transition is genuinely a boundary, not a gradual fade: the run-to-run spread of outcomes, which is large exactly where a run can tip either way, spikes in a narrow band around *inter* $\approx 10^{-3}$ and is flat on both sides. Why is contact so corrosive? Because fusion shares *evidence*, not just conclusions: an open vanguard of just 5% of the population lifts end-state conversion from 0.60 to 0.79, because the vanguard’s anomalous data flows across the trust channels and sets off each dogmatist’s *own* crisis. An agent can silence its own instruments; it cannot silence its neighbours’. Gating cannot defeat connection.

A control pins down what does the stalling. Keep conviction exactly as high but switch the gate off—the agents stay just as motivated, but keep listening—and the community revolts in full (revolted fraction 1.0, versus 0.0 once the gate silences the disconfirming channels, $\gamma \rightarrow 1$). So the evidence was never absent or out of reach: what stalls a community is conviction closing the channels that would deliver it. Lock-in is derived, not assumed—it appears exactly when enough disconfirming precision is silenced (Eq. 4) that the data lose their grip on the agent’s commitments.

Discovering the unconceived is cheap; keeping it requires discovering together. Discovery in this model means waking a node—oxygen—that no agent’s network yet contains, and since one cannot search for a concept one has not conceived, candidates arrive at random (Poisson rate r , §3.3). Individually this behaves exactly as a waiting process should: the fraction of the open community that discovers within the horizon follows the textbook law $1 - (1 - r)^T$, rising from 12% at $r = 0.01$ to 88% at $r = 0.16$ (Fig. 5, left); and because ~ 40 agents draw in parallel, *somebody* wakes the node within a few steps at every rate (right panel)—the community’s first discovery is




Figure 4: Only isolation protects a superseded paradigm (phlogiston sweep). *Left, centre:* the fraction of the dogmatic community that converts to oxygen, as a function of contact with the open community (*inter*) and of how strongly conviction gates disconfirming channels (*s*), at two conviction levels ($\lambda_d = 0.35$ and 0.8). Gating keeps conversion low only when contact is essentially zero; past $inter \approx 0.005$ nearly everyone converts, and doubling conviction barely changes the picture. *Right:* the run-to-run spread of outcomes—large exactly where a run can tip either way—spikes in a narrow band around $inter \approx 5 \times 10^{-4}$, the signature of a sharp boundary rather than a gradual fade.

nearly free. What is not free is survival. A discovery made by one agent at a time *never* survives, at any rate (the red curve): the moment the lone discoverer fuses beliefs with peers who lack the concept, its oxygen node is averaged away—the shared prior treats the slot the others do not represent as a firmly held “no such thing,” and pulls the awakened coupling back to zero. Only when the proposal wakes the community near-simultaneously does the concept take root (47% of runs under simultaneous proposal). A new concept does not need a critical mass of believers eventually; it needs them at the same *time*. (One modeling caveat: under posterior fusion the unconceived slot’s pin is pooled as if it were a held belief. This reads naturally as incommensurability—the shared conceptual scheme suppresses concepts it does not contain—but fusion schemes that exclude unconceived dimensions would soften the effect.)

The cycle is robust; the one fragile dial is how loudly the world must knock. Re-running the full arc across forgetting and observation noise ($\omega \in [0.9, 1.0]$, $\sigma_o \in [0.25, 1.0]$) changes nothing: crisis, reduction and discovery all still arrive, even in the no-forgetting corner $\omega = 1$. The one dial that can break the cycle is the expansion trigger—how large an unexplained residual an agent demands before it will even entertain a new node. Set at 1.0–1.5 the cycle completes everywhere; at 2.5 the bar sits *above* the prediction error the dead paradigm actually leaves behind after reduction, so the new node is never proposed and discovery dies almost everywhere (fraction ≤ 0.42 , mostly 0). The calibration lesson is simple: an agent must be willing to consider new structure at the level of error its old structure already produces.

4.2 Limitations and future work

One prediction fails cleanly, and it names the next modelling target. The structural account predicts a belt-first, core-last *staircase*: the cheap periphery reorganizes first, the costly core follows once enough evidence accumulates to pay its carry-over. A conservatism-gated revision rate does *not* recover it: the high- κ nodes reorganize *first*, and the κ -gradient gate is indistinguishable from uniform and shuffled- κ nulls (Appendix E gives the full analysis). The root cause is that the Fisher deposit is set by the observation operator rather than the data, so theory identity rides the posterior




Figure 5: Discovery and survival of the unconceived (phlogiston sweep). *Left*: the fraction of the open community that discovers oxygen follows the waiting-time law $1 - (1 - r)^T$ (dashed), but the fraction whose concept *survives* fusion is zero at every Poisson rate (red): staggered discoverers are pulled back by the communal prior, while simultaneous proposal (point, right edge) survives in 47% of runs—critical mass in *time*. *Right*: with ~ 40 agents drawing in parallel, the population’s first discovery is nearly free at any rate.

means and κ , though real and computable (Appendix B), never enters the update. The first item of future work is therefore a *data-dependent* observation model in which the couplings themselves are learned, making conservatism load-bearing in the dynamics. The second is the utility source u , here exogenous and static: letting it accrete onto structurally entrenched commitments— u aligning with κ as a programme matures—would recover Lakatos’s progressive and degenerating problemshifts and meet Laudan’s reticulated co-variation of theory, method, and aim (Lakatos, 1970; Laudan, 1984).

The model is, finally, stylized: the nets are illustrative, their weights chosen to expose the geometry; the dials of Table 1 are swept, not inferred; and a community that “tracks” the eventually-true theory is committed to it, not thereby correct. What clears the “more than a redescription” bar is not any single result but the conjunction a scalar credence cannot voice: the Schur-complement degeneracy (a scalar has no off-diagonal to fill), the derived-not-imposed lock-in with its deaf-but-honest control, and the expansion that wakes an unconceived node only when the world makes its absence felt.

5 Discussion

Everything distinctive follows from taking the hidden structure seriously: a conservatism gradient derived rather than assigned, an exploration drive that is the value of expanding the model rather than a tuned bonus, a conviction field that is the same operator on a different source, and a population arc—crisis, revolution, lock-in, or release—that emerges as threshold crossings of the agents’ own evidence. The sweep adds the population-scale phenomenology: the lock-in boundary lives at very weak coupling, so the protective strategies internal to a community (conviction, gating) are dominated by the single external variable of contact; and the survival of a new concept is governed not by how many eventually conceive it but by how nearly in time they do.

We should be candid about which philosopher this picture serves: Lakatos, more than Kuhn. The

conservatism geometry *derives* his hard core and protective belt—a node is costly to revise because much depends on it—rather than decreeing them (Lakatos, 1970). Identifying the carry-over cascade with Kuhnian *incommensurability* is the framing’s most contestable move: a directed Bayes net makes incommensurability commensurable by construction—one precision space, revolution a long traverse rather than a gestalt switch—whereas Kuhn’s incommensurability is precisely the absence of the common measure the network presupposes (Kuhn, 1983). The staircase null (§4.2) is the tell: graded reorganization on fixed wiring is a Lakatosian problemshift, not a Kuhnian reconstitution. Yet the sweep’s survival result earns a piece of it back: the fusion that crushes a staggered discovery is a shared conceptual scheme suppressing a concept it does not contain—incommensurability as a population-level force, not a stipulated incomparability.

Finally, the unconceived marks the limit of structure learning itself. Structure learning navigates beautifully within the dimensions a model posits but is blind to the dimension no one has proposed—Stanford’s problem of unconceived alternatives (Stanford, 2006): not uncertainty over which known structure is right, but whether the space of known structures contains the truth at all. One cannot place a posterior over the unconceived; running an expansion process that keeps proposing nodes at some Poisson rate is what a healthy epistemic system does instead of pretending it can. §4.1 shows both edges of that blade: too low a rate leaves the dimension undiscovered within the horizon, and a discovery made alone does not survive the communal prior. The hidden structure of the world, which forced us into structure learning on the first page, is in the end the thing no amount of structure learning can guarantee to have caught.

References

- M. Albarracin, D. Demekas, M. J. D. Ramstead, and C. Heins. Epistemic communities under active inference. *Entropy*, 24(4):476, 2022.
- J. Dalege, D. Borsboom, F. van Harreveld, H. van den Berg, M. Conner, and H. L. J. van der Maas. Toward a formalized account of attitudes: the Causal Attitude Network (CAN) model. *Psychological Review*, 123(1):2–22, 2016.
- J. Friedman, T. Hastie, and R. Tibshirani. Sparse inverse covariance estimation with the graphical lasso. *Biostatistics*, 9(3):432–441, 2008.
- K. J. Friston and W. Penny. Post hoc Bayesian model selection. *NeuroImage*, 56(4):2089–2099, 2011.
- K. Friston, F. Rigoli, D. Ognibene, C. Mathys, T. Fitzgerald, and G. Pezzulo. Active inference and epistemic value. *Cognitive Neuroscience*, 6(4):187–214, 2015.
- K. J. Friston, M. Lin, C. D. Frith, G. Pezzulo, J. A. Hobson, and S. Ondobaka. Active inference, curiosity and insight. *Neural Computation*, 29(10):2633–2683, 2017.
- K. J. Friston, T. Parr, and P. Zeidman. Bayesian model reduction. arXiv:1805.07092, 2018.
- D. Hyland and M. Albarracin. On the variational costs of changing our minds. arXiv:2509.17957, 2025. (Also in *Active Inference: IWAJ 2025*, CCIS vol. 2857, Springer, 2026, doi:10.1007/978-3-032-16955-6_1.)
- N. Jonard, S. Reijula, and L. Marengo. Theory choice in epistemic networks: five ways to avoid premature convergence. PhilSci Archive, 2025.
- P. Kitcher. The division of cognitive labor. *Journal of Philosophy*, 87(1):5–22, 1990.
- D. Koller and N. Friedman. *Probabilistic Graphical Models: Principles and Techniques*. MIT Press, 2009.
- T. S. Kuhn. *The Structure of Scientific Revolutions*. University of Chicago Press, 1962.
- D. J. C. MacKay. Bayesian interpolation. *Neural Computation*, 4(3):415–447, 1992.
- B. Millidge, A. Tschantz, and C. L. Buckley. Whence the expected free energy? *Neural Computation*, 33(2):447–482, 2021.
- R. M. Neal. *Bayesian Learning for Neural Networks*. Lecture Notes in Statistics 118, Springer, 1996.
- C. T. Nguyen. Echo chambers and epistemic bubbles. *Episteme*, 17(2):141–161, 2020.
- C. O’Connor and J. O. Weatherall. *The Misinformation Age*. Yale University Press, 2019.
- R. Smith, P. Schwartenbeck, T. Parr, and K. J. Friston. An active inference approach to modeling structure learning: concept learning as an example case. *Frontiers in Computational Neuroscience*, 14:41, 2020.
- P. K. Stanford. *Exceeding Our Grasp: Science, History, and the Problem of Unconceived Alternatives*. Oxford University Press, 2006.
- M. Strevens. The role of the priority rule in science. *Journal of Philosophy*, 100(2):55–79, 2003.
- P. Thagard. Explanatory coherence. *Behavioral and Brain Sciences*, 12(3):435–467, 1989.
- M. E. Tipping. Sparse Bayesian learning and the relevance vector machine. *Journal of Machine Learning Research*, 1:211–244, 2001.

- K. J. S. Zollman. The epistemic benefit of transient diversity. *Erkenntnis*, 72(1):17–35, 2010.
- T. L. Griffiths and Z. Ghahramani. The Indian Buffet Process: an introduction and review. *Journal of Machine Learning Research*, 12:1185–1224, 2011.
- C. Kemp and J. B. Tenenbaum. The discovery of structural form. *Proceedings of the National Academy of Sciences*, 105(31):10687–10692, 2008.
- T. D. Ullman, N. D. Goodman, and J. B. Tenenbaum. Theory learning as stochastic search in a language of thought. *Cognitive Development*, 27(4):455–480, 2012.
- K. Friston, T. FitzGerald, F. Rigoli, P. Schwartenbeck, and G. Pezzulo. Active inference: a process theory. *Neural Computation*, 29(1):1–49, 2017.
- C. Mathys, J. Daunizeau, K. J. Friston, and K. E. Stephan. A Bayesian foundation for individual learning under uncertainty. *Frontiers in Human Neuroscience*, 5:39, 2011.
- R. Smith, K. J. Friston, and C. Whyte. A step-by-step tutorial on active inference and its application to empirical data. *Journal of Mathematical Psychology*, 107:102632, 2022.
- L. Da Costa, T. Parr, N. Sajid, S. Veselic, V. Neacsu, and K. Friston. Active inference on discrete state-spaces: a synthesis. *Journal of Mathematical Psychology*, 99:102447, 2020.
- F. L. Holmes. The revolution in chemistry and physics: overthrow of a reigning paradigm or competition between contemporary research programs? *Isis*, 91(4):735–753, 2000.
- P. Thagard. The conceptual structure of the chemical revolution. *Philosophy of Science*, 57(2):183–209, 1990.
- V. D. Boantz and O. Gal. The ‘absolute existence’ of phlogiston: the losing party’s point of view. *British Journal for the History of Science*, 44(3):317–342, 2011.
- G. Blumenthal and J. Ladyman. The development of problems within the phlogiston theories, 1766–1791. *Foundations of Chemistry*, 19(3):241–280, 2017.
- B. Ball, A. Kolioussis, A. Mohanan, and M. Peacey. Computational philosophy: reflections on the PolyGraphs project. *Humanities and Social Sciences Communications*, 11:186, 2024.
- M. Michelini, J. Osorio, W. Houkes, D. Šešelja, and C. Straßer. Scientific disagreements and the diagnosticity of evidence: how too much data may lead to polarization. *Journal of Artificial Societies and Social Simulation*, 26(4):5, 2023.
- M. Albarracin, R. J. Pitliya, T. St. Clere Smithe, D. A. Friedman, K. Friston, and M. J. D. Ramstead. Shared protentions in multi-agent active inference. *Entropy*, 26(4):303, 2024.
- K. J. Friston, T. Parr, C. Heins, A. Constant, D. Friedman, T. Isomura, C. Fields, T. Verbelen, M. Ramstead, J. Clippinger, and C. D. Frith. Federated inference and belief sharing. *Neuroscience & Biobehavioral Reviews*, 156:105500, 2024.
- P. T. Waade, C. L. Olesen, J. E. Laursen, S. W. Nehrer, C. Heins, K. Friston, and C. Mathys. As one and many: relating individual and emergent group-level generative models in active inference. *Entropy*, 27(2):143, 2025.
- T. D. Ullman and J. B. Tenenbaum. Bayesian models of conceptual development: learning as building models of the world. *Annual Review of Developmental Psychology*, 2:533–558, 2020.
- R. Smith, V. Neacsu, M. B. Mirza, R. A. Adams, and K. J. Friston. Simulating empirical phenomena of inference under active learning and active selection through Bayesian model reduction. *PLOS ONE*, 17(12):e0277199, 2022.
- M. Priorelli and I. P. Stoianov. Dynamic inference by model reduction. *bioRxiv*, 2023. doi:10.1101/2023.09.10.557043.
- K. J. Friston, V. Litvak, A. Oswal, A. Razi, K. E. Stephan, B. C. M. van Wijk, G. Ziegler, and P. Zeidman. Bayesian model reduction and empirical Bayes for group (DCM) studies. *NeuroImage*, 128:413–431, 2016.
- K. J. Friston et al. Designing ecosystems of intelligence from first principles. *Collective Intelligence*, 3(4), 2024. arXiv:2212.01354.
- R. Kaufmann, P. Gupta, and J. Taylor. An active inference model of collective intelligence. *Entropy*, 23(7):830, 2021.
- D. Acemoglu and A. Ozdaglar. Opinion dynamics and learning in social networks. *Dynamic Games and Applications*, 1(1):3–49, 2011.
- V. Bala and S. Goyal. Learning from neighbours. *Review of Economic Studies*, 65(3):595–621, 1998.
- J. A. Bohren and D. N. Hauser. Learning with heterogeneous misspecified models: characterization and robustness. *Econometrica*, 89(6):3025–3077, 2021.
- P. Feyerabend. *Against Method*. New Left Books, 1975.
- C. Genest and J. V. Zidek. Combining probability distributions: a critique and an annotated bibliography. *Statistical Science*, 1(1):114–135, 1986.
- R. Hegselmann and U. Krause. Opinion dynamics and bounded confidence: models, analysis and simulation. *Journal of Artificial Societies and Social Simulation*, 5(3), 2002.
- T. S. Kuhn. Commensurability, comparability, communicability. *PSA: Proceedings of the Biennial Meeting of the Philosophy of Science Association*, 1982(2):669–688, 1983.
- I. Lakatos. Falsification and the methodology of scientific research programmes. In I. Lakatos and A. Musgrave, editors, *Criticism and the Growth of Knowledge*, pages 91–196. Cambridge University Press, 1970.
- L. Laudan. *Progress and Its Problems: Towards a Theory of Scientific Growth*. University of California Press, 1977.
- L. Laudan. *Science and Values: The Aims of Science*

- and Their Role in Scientific Debate*. University of California Press, 1984.
- K. Lehrer and C. Wagner. *Rational Consensus in Science and Society*. D. Reidel, 1981.
- C. O'Connor and J. O. Weatherall. Scientific polarization. *European Journal for Philosophy of Science*, 8(3):855–875, 2018.
- K. J. S. Zollman. Network epistemology: communication in epistemic communities. *Philosophy Compass*, 8(1):15–27, 2013.

A The operator, the tilt, and the Schur residue

Linear-Gaussian construction. The couplings form a structural equation model $x = Ax + \zeta$, $\zeta \sim \mathcal{N}(0, D^{-1})$ with A the weighted DAG adjacency (Koller and Friedman, 2009); the choice is for tractability, and a nonlinear coupling would replace T by a local linearization. The propagation operator is the resolvent

$$\mathsf{T} = (I - A)^{-1} = \sum_{k \geq 0} A^k, \quad \mathsf{T}_{iw} = \sum_{\text{paths } i \rightsquigarrow w} \prod \pi = \pi_{i \rightsquigarrow w}, \quad (6)$$

the Neumann series finite because a DAG makes A nilpotent. We realise T on the coupling magnitudes $|A|$ so downstream mass does not cancel when a child opposes a parent; the signed A builds the precision matrix. We index T so that κ_i sums row i ; under the SEM convention the released code stores the transpose and computes the same fields as column-sums of $\mathsf{T}^\top$, so a reader reproducing $\mathsf{T}\mathbf{1}$ literally against the code’s array must transpose first.

The value tilt. With honest posterior $p(x \mid o, G) = \mathcal{N}(\mu, \Sigma)$, the tilted posterior of Eq. (3) is a pure mean shift, $q_\lambda(x) = \mathcal{N}(\mu + \lambda \Sigma U, \Sigma)$, with cumulant-generating function $Z(\lambda) = \lambda U^\top \mu + \frac{1}{2} \lambda^2 U^\top \Sigma U$. Note that gating a channel ($\rho_k \rightarrow 0$) *raises* the posterior variance and hence Z'' ; lock-in lives in the vanishing mean-update gain $\rho_k H_k \Sigma U$, not in the curvature.

The Schur residue. Marginalizing a hub condenses its influence into new edges among its neighbours: for a common-cause structure over (x_1, x_2, b) ,

$$\Pi = \begin{pmatrix} 2 & 0 & 1 \\ 0 & 2 & 1 \\ 1 & 1 & 2 \end{pmatrix} \xrightarrow{\text{marginalize } b} \begin{pmatrix} 2 & 0 \\ 0 & 2 \end{pmatrix} - \frac{1}{2} \begin{pmatrix} 1 \\ 1 \end{pmatrix} \begin{pmatrix} 1 & 1 \end{pmatrix} = \begin{pmatrix} 1.5 & -0.5 \\ -0.5 & 1.5 \end{pmatrix}, \quad (7)$$

a direct coupling induced where the joint had a zero. Conditioning on a shared *effect* (the supernova collider of Appendix B) is the dual operation, and both fill in the off-diagonal coupling the carry-over cost is built from.

B The two fields on a worked example: the dark-energy net

This appendix realises the two fields of Fig. 1 on a domain net: general relativity (GR) at the root; the density parameters Ω_m and the cosmological constant Λ as dials; the Friedmann expansion, distance modulus $\mu(z)$ and a Type-Ia supernova (SN) reading as a descending chain; light-bending, the perihelion precession, gravitational redshift, CMB and BBN as GR’s other branches; and a modified-gravity (MOND) prediction the community wants false. The net is hand-wired to expose the geometry, not fitted to data.

Computing $\kappa = \mathsf{T}\mathbf{1}$ and $U = \mathsf{T}u$ on this net (Fig. 6) makes the two-field claim concrete. Conservatism is steepest at GR ($\kappa_{\text{GR}} = 10.5$, the maximum) and shallow at the cosmological constant ($\kappa_\Lambda = 2.85$, a ratio of 3.7): GR carries every predicted observable below it, Λ only the distance–redshift branch. This is Lakatos’s belt-protects-core falling straight out of the operator—when the 1998 anomaly arrives at the bottom, the cheap move is to turn the dial with the smallest descendant set (Λ), not the ancestor with the largest (GR). The conviction field $U = \mathsf{T}u$, with the community valuing the accelerating-universe supernova reading ($u_{\text{SN}} = +1$), the Λ commitment (+0.6) and wanting MOND false (−0.8), *inverts* that order: $U_\Lambda = 1.19$ (the most cherished) exceeds $U_{\text{GR}} = 0.45$ while $\kappa_\Lambda < \kappa_{\text{GR}}$. The two fields are de-correlated (Pearson $r(\kappa, U) = 0.25$), the

visible face of the algebraic independence guaranteed whenever $u \not\propto \mathbf{1}$: Λ is *cheap yet cherished*, GR *central yet comparatively value-neutral*, the GR machinery (light-bending, perihelion, clocks, CMB, BBN) central yet value-neutral ($U \approx 0$), and MOND the conclusion the community wants false ($U < 0$)—exactly the three annotated cases of Fig. 1, now computed by the operator on a domain net rather than assigned by hand.




Figure 6: The two fields, computed on the dark-energy paradigm net. *Left*: conservatism $\kappa = T\mathbf{1}$ (Eq. 1)—GR is hot ($\kappa = 10.5$), the observation leaves cold; the directed arrows make “everything hangs below GR” visible. *Right*: conviction $U = T u$ (Eq. 2) on the same graph— Λ and the supernova reading green (cherished), MOND red (wanted false), GR’s other machinery value-neutral. The decoupling is the realisation of Fig. 1: the value order inverts the cost order ($U_\Lambda > U_{GR}$ while $\kappa_\Lambda < \kappa_{GR}$; $r(\kappa, U) = 0.25$).

The Schur residue of Eq. (7) is equally concrete. A priori Ω_m and Λ are independent (correlation 0.00, a round cloud). Conditioning on the shared supernova observation—clamping the descendant both dials feed—manufactures an off-diagonal coupling between them where the joint had a zero: the cloud collapses onto the tilted “banana” (Fig. 7), with correlation rising to +0.50. This is the Ω_m – Λ degeneracy direction every dark-energy contour plot draws, here reproduced as a Schur complement on the precision matrix. Conditioning on a shared *effect* (the supernova collider—explaining-away) and marginalizing a shared *cause* (Eq. 7) are *dual* operations, not the same one, but both fill in the off-diagonal coupling the carry-over cost is built from.

C Coupling of the two fields under model expansion

The two fields are coupled by the expansion operation, and with T defined the coupling follows directly. When the agent expands its model—proposing a new commitment N wired by candidate conditionals to a set of existing nodes—the joint gains an explicit factor,

$$p(x, x_N \mid G') \propto p(x \mid G) \prod_{(N \leftrightarrow v)} \psi_{Nv}(x_v, x_N)^{\pi_{Nv}}, \quad (8)$$

with the candidate edges entered at free prior precision ($\alpha_e \rightarrow 0$) so that reduction can carve them back. In operator terms, adding N appends a row and a column to A —and therefore to T —and appends one entry to *each* of the two source vectors, the unit source $\mathbf{1}$ and the utility source u . The




Figure 7: The Schur degeneracy, computed. The two dials are uncorrelated a priori (grey, $r = 0.00$); conditioning on their shared supernova descendant induces the off-diagonal coupling (crimson, $r = +0.50$)—the Ω_m - Λ “banana”. Marginalizing a shared latent hub and conditioning on a shared observation are the two readings of one complement.

two consequences are then immediate. Every node N touches gains the new coupling precision it must now also propagate,

$$\Delta\kappa_v = \pi_{Nv}, \quad (9)$$


while the conviction field gains, at N and along its new edges,

$$\Delta U = u_N + \sum_{w \in \text{desc}(N)} \pi_{N \sim w} u_w. \quad (10)$$

A single expansion thus writes the new commitment into both the cost geometry and the value geometry (Fig. 8). This is the precise content of motivated structure learning: an agent contemplating a node that would, once wired in, lend value to a cherished core—or force the revision of one—registers that prospect in *both* fields, as a carry-over cost it must pay and as a value change it may welcome or resist, and scores the keep-or-prune decision against the value-tilted ledger of Eq. (3). The same proposal can therefore be evidence-favoured and value-rejected, or the reverse.

D A coarse model resolves a complex world only slowly

This experiment isolates the cost of a coarse Bayes net against a world richer than it. The true world is a complex twelve-parameter Gaussian with four latent factors (correlated, individually ambiguous directions); a coarse agent senses only res random aggregate combinations per step. Coarser is slower (Fig. 9, left): the time to resolve the world state to within a fixed tolerance grows from 58 steps at full resolution to 221 steps at $\text{res} = 1$, and below a resolution the agent does not finish within the horizon. Adding forgetting floors the resolution outright (Fig. 9, right): a net that both under-senses and forgets plateaus at a residual error that rises as ω falls ($0.08 \rightarrow 0.28$), and at $\text{res} = 3$, $\omega \leq 0.93$



(a) enters the belief network (coupling π , $\Delta\kappa_v = \pi_{Nv}$) (b) enters the utility network (propagated value, ΔU)

Figure 8: One expansion, two fields. Adding a commitment N with new conditionals appends a row and a column to the operator $\mathbf{T}$ and an entry to each source vector: into the belief network, where each new edge carries a coupling precision π and so adds to carry-over cost by Eq. (9) (a), and into the utility network, where N 's value propagates along the same new edges by Eq. (10) and tints the commitments it reaches (b). The keep-or-prune decision is scored against the value-tilted ledger, so a proposal can be evidence-favoured yet value-rejected, or the reverse.

never resolves the world at all. A complex world has no clear answer for a bounded, forgetful, under-sensing learner—the honest limit the hidden-structure premise of §2 implies.




Figure 9: A coarse net on a complex world. a Relative error in the reconstructed world state over time, by sensing resolution; time-to-resolve grows sharply as the net coarsens. b With a fixed coarse sensor, forgetting plateaus the error at a floor that rises as ω falls—the world is never fully resolved.

E The staircase negative in detail

In the base dynamics the conservatism κ is a static read-out that never enters the update. The natural fix is to let κ gate the *rate* of revision: a per-node deposit gate $g_i = 1/(1 + \beta\kappa_i)$, PD-preserving, so a high- κ core node integrates evidence slower than a cheap belt node. We tested it on the phlogiston net, and it does *not* recover the staircase (Fig. 10). The belt $\rightarrow$ core lag stays negative (-55 steps at $\beta = 0$, -45 at large β): the high- κ nodes reorganize *first*, not last, and the κ -gradient gate is indistinguishable from a gradient-free uniform slowdown (-45) and a shuffled- κ gate (-45). Two reasons, which share one root cause. First, a deposit gate scales both numerator and denominator of the mean $\mu = \Pi^{-1}h$, so it barely throttles mean-tracking—and mean-tracking is

where theory identity lives. Second, the nodes that actually reorganize are exactly the high- κ mass-law commitments the relational operator refutes directly, while the low- κ belt are agreement nodes with nothing to reorganize, so the conservatism stratification is confounded with move-versus-don't-move. The root cause: with an operator-set Fisher deposit $J = H^\top H / \sigma^2$ —set by the observation operator, not the data—the structure does not move, so theory identity rides the posterior *means* and neither a rate gate on those means nor the stratified couplings can express the staircase. The negative localizes the missing ingredient to a data-dependent observation model in which the couplings themselves are learned (a structure-matched or nonlinear operator).




Figure 10: The staircase is *not* recovered by a conservatism-gated revision rate (phlogiston). *Left*: reorganization by κ -shell at the best β ($\beta = 0$ baseline dotted)—the high- κ “core” moves first and the low- κ “belt” (agreement nodes) never reorganizes, the reverse of the prediction. *Right*: the belt→core lag stays negative and the κ -gradient gate (purple) is matched by the gradient-free uniform (grey square) and shuffled- κ (red cross) nulls—the gradient adds nothing. The negative localizes the missing ingredient to data-dependent structure learning.